

a) Ans:


```
> n = 100
> X = rnorm(n)
> Epsilon = rnorm(n)
> |
```

b) Ans:

```
> beta0 = 1; beta1 = 2; beta2 = 3; beta3 = 4
> Y <- beta0 + beta1*X + beta2*X^2 + beta3*X^3 + Epsilon
> Y
 [1]  9.11268604 -0.50328770 46.64200818  0.05754668 30.20492616 24.34156438
 [7]  0.82223478  3.53397272 -0.76005025  1.39627553 10.51294417  1.88403350
[13]  1.46057424 -0.92160051  0.42375478 11.97992270  1.29905042  0.54365751
[19] -6.66357258  4.91391408 19.13191054 23.42737098  6.71856443 -20.15866600
[25]  3.67884224  4.52841080 -1.14595307  0.66843265 -1.70978672  4.31485803
[31] -1.50649519 -9.70358411 -1.06009237 19.44949972  1.11023830  5.45179284
[37]  9.86590802  0.98314072 -0.15484853  0.83421818  5.30452060 -0.77231186
[43]  1.71572230  1.42819428 -34.40847174 -10.05665214  7.05769676  0.27510897
[49]  6.85248060 40.01686832 24.23343629  4.36882129  1.65958637 -0.85836097
[55] 25.88529455  4.16311864 -3.17132548 -1.35157575 -21.05889024  1.22985431
[61] -54.16065954 20.04821140  1.15958422  0.97220471  0.73283733  4.54609506
[67]  4.27975086  4.46444523  1.56781446 -5.20251447 -0.37232527  0.25452041
[73] -0.95385683  1.31611976 -2.26215292 -0.33502269  1.68046445  0.99317092
[79]  1.92911870 35.76894388  5.68638145 11.27359527 -1.95433921  1.06533896
[85] 15.82146580  3.34235027 -28.44248461  1.16681418 -21.54751035 -0.26866779
[91] 17.10024919  6.58628618 11.07689075  2.12162723  4.62101693 -2.55323488
[97] 28.27279003 -0.01043066 -0.01126861  6.69387686
> |
```

c) Ans:

```
> library(leaps)
> data = data.frame(Y = Y, X1 = X, X2 = X^2, X3 = X^3, X4 = X^4, X5 = X^5, X6 = X^6, X7 = X^7, X8 = X^8, X9 = X^9, X10 = X^10)
> best_subset = regsubsets(Y ~ ., data = data, nvmax = 10)
> summary.reg = summary(best_subset)
> summary.reg$cp
 [1] 1086.104639  85.498864  6.186395  7.383294  6.092477  5.288440  6.128827
 [8]  7.002717  9.000439 11.000000
> summary.reg$bic
 [1] -231.1328 -415.0159 -470.7408 -466.9570 -465.7894 -464.2097 -460.8829 -457.5350 -452.9324
[10] -448.3278
> summary.reg$adjr2
 [1] 0.9086702 0.9859889 0.9922560 0.9922385 0.9924210 0.9925678 0.9925825 0.9925947 0.9925126
[10] 0.9924285
> par(mfrow = c(1,3))
> plot(summary.reg$cp, xlab = "Number of Variables", ylab = "Cp", type = "b")
> plot(summary.reg$bic, xlab = "Number of Variables", ylab = "BIC", type = "b")
> plot(summary.reg$adjr2, xlab = "Number of Variables", ylab = "Adjusted R2", type = "b")
> coef(best_subset, 3)
(Intercept)      X1      X2      X3
 1.113556  2.059186  2.836236  3.974698
> coef(best_subset, 4)
(Intercept)      X1      X2      X3      X4
 1.13667646  1.85868682  2.79845558  4.16948947 -0.03207058
> |
```



From the output we know that the model with 3 variables (X , X^2 , X^3) is the best choice.

- We choose the model with the lowest BIC (because lower BIC = better model with good fit and less complexity). Best BIC is at 3 variables (-470.74) because it's the smallest (most negative) value.
- We want Adjusted R^2 to be as high as possible, but we also look for where it stops improving a lot. Adjusted R^2 improves a lot up to 3 variables, but then flattens after that ($0.992 \rightarrow 0.992 \rightarrow 0.992$). This means adding more variables doesn't make the model better after 3 variables.
- Here we choose a model with C_p close to the number of variables. Best C_p is at 3 or 4 variables, but again, 3 is enough because it's closer and matches BIC.
- Coefficients for the 3-variable model match the true cubic relationship we simulated (X , X^2 , X^3).

Best Model:


Hence, the model with 3 variables (X , X^2 , X^3) is the best choice.


d) Ans:

```
> forward_model = regsubsets(Y ~ ., data = data, nvmax = 10, method = "forward")
> summary.forward = summary(forward_model)
> par(mfrow = c(1, 3))
> plot(summary.forward$cp, xlab = "Number of Variables", ylab = "forward_Cp", type = "b")
> plot(summary.forward$bic, xlab = "Number of Variables", ylab = "forward_BIC", type = "b")
> plot(summary.forward$adjr2, xlab = "Number of Variables", ylab = "forward_Adjusted R2", type = "b")
> backward_model = regsubsets(Y ~ ., data = data, nvmax = 10, method = "backward")
> summary.backward = summary(backward_model)
> par(mfrow = c(1, 3))
> plot(summary.backward$cp, xlab = "Number of Variables", ylab = "backward_Cp", type = "b")
> plot(summary.backward$bic, xlab = "Number of Variables", ylab = "backward_BIC", type = "b")
> plot(summary.backward$adjr2, xlab = "Number of Variables", ylab = "backward_Adjusted R2", type = "b")
> summary.forward$cp
[1] 1086.104639 85.498864 6.186395 7.383294 6.092477 7.961659 7.600343
[8] 8.419645 10.349037 11.000000
> summary.backward$cp
[1] 2618.160625 1159.991107 331.530656 59.990047 14.423327 11.895859 6.128827
[8] 7.002717 9.000439 11.000000
> coef(forward_model, 3)
(Intercept) X1 X2 X3
1.113556 2.059186 2.836236 3.974698
> coef(backward_model, 4)
(Intercept) X1 X2 X5 X7
1.0821532 4.6766938 2.9848434 1.3672846 -0.1249905
>
```

Plot Zoom

— □ ×





Forward Stepwise Selection

The results are exactly the same as Best Subset Selection in 8(c). Cp, BIC, and Adjusted R² all favor the 3-variable model. Hence, Selected model is 3 variables (X , X^2 , X^3) with the same coefficients as above in (c).

Backward Stepwise Selection

Cp is slightly lower for 6–7 variables, but differences are small. BIC still prefers smaller models (around 3 variables). And adjusted R² increases slightly for larger models but plateaus after 3 variables. Hence, backward may include extra variables, like X_5 or X_7 , but does not significantly improve fit.


e) Ans:


Using the optimal λ (left dotted line), Lasso selected 3 predictors out of the 10 possible. The minimum error (left line) is at a $\log(\lambda)$ around -3. At this λ , the model keeps 3 non-zero predictors. Hence Lasso identifies the cubic relationship and removes noise variables.

```

> x.mat = model.matrix(Y ~ poly(X, 10, raw = TRUE))[, -1]
> y.vec = Y
> set.seed(1)
> cv.lasso = cv.glmnet(x.mat, y.vec, alpha = 1)
> plot(cv.lasso)
> best.lambda = cv.lasso$lambda.min
> best.lambda
[1] 0.04757254
> lasso.fit = glmnet(x.mat, y.vec, alpha = 1, lambda = best.lambda)
> coef(lasso.fit)
11 x 1 sparse Matrix of class "dgCMatrix"
      s0
(Intercept) 1.172886
poly(X, 10, raw = TRUE)1 2.089356
poly(X, 10, raw = TRUE)2 2.768111
poly(X, 10, raw = TRUE)3 3.940270
poly(X, 10, raw = TRUE)4 .
poly(X, 10, raw = TRUE)5 .
poly(X, 10, raw = TRUE)6 .
poly(X, 10, raw = TRUE)7 .
poly(X, 10, raw = TRUE)8 .
poly(X, 10, raw = TRUE)9 .
poly(X, 10, raw = TRUE)10 .
>


```

 Plot Zoom



f) Ans:

```
> beta=8
> Y_new = beta + beta*X^7 + Epsilon
> Y_new
[1] 7.672349e-02 1.042172e+00 -2.186978e+00 2.115116e+02 3.487894e-01 7.650159e-01 1.769004e+00 2.866976e+00
[9] 1.552024e+00 2.680194e+00 1.447461e+02 5.493029e-01 2.146583e+00 -2.090363e+03 1.903036e+01 6.071921e-01
[17] 6.800071e-01 6.058738e+00 3.509355e+00 1.031162e+00 4.922207e+00 3.775448e+00 7.854207e-01 -9.856202e+02
[25] 1.180985e+00 1.712666e+00 9.264178e-01 -1.181253e+02 2.726265e-01 6.935494e-01 6.943704e+01 4.111045e-01
[33] 1.542024e+00 -5.183941e-01 -7.381404e+01 -5.534083e-01 6.871716e-01 4.717201e-01 1.594016e+01 2.149422e+00
[41] -9.143855e-01 2.176047e+00 -2.588587e-02 6.689762e-01 -7.041491e-01 -4.606489e-01 3.094016e+00 2.284252e+00
[49] -2.863023e-01 2.657716e+00 1.462866e+00 7.241157e-01 6.862313e-01 -1.867608e+01 9.879296e+01 9.557099e+02
[57] 1.992825e+00 -1.044521e+01 -2.285727e-01 2.869284e+00 3.687950e+03 7.613529e-01 2.652618e+00 1.886423e+00
[65] -6.218468e-01 3.206171e+00 -4.985541e+02 1.157483e+02 8.556163e-01 1.829159e+03 3.351953e+00 3.785770e-01
[73] 1.710521e+00 -4.041200e+00 -3.826395e+01 9.667029e-01 1.760729e+00 3.075245e+00 2.027393e+00 2.009945e+00
[81] -3.851762e-01 1.983889e+00 2.641592e+01 -1.529148e+02 1.729625e+00 8.448741e-01 1.474204e+01 2.319903e-01
[89] 5.773855e-01 7.466638e-02 7.122336e-01 3.140924e+01 2.293300e+01 2.490616e+00 2.024712e+02 8.759019e-02
[97] -4.176221e+01 -1.786196e-01 -3.163115e+01 5.762963e-01
> best_subset_new = regsubsets(Y ~ ., data , nvmax = 10)
> summary_new = summary(best_subset_new)
> par(mfrow = c(1, 3))
> plot(summary_new$cp, xlab = "Number of Variables", ylab = "Cp", type = "b")
> plot(summary_new$bic, xlab = "Number of Variables", ylab = "BIC", type = "b")
> plot(summary_new$adjr2, xlab = "Number of Variables", ylab = "Adjusted R2", type = "b")
> coef(best_subset_new, 7)
(Intercept)          X1          X2          X3          X5          X6          X8          X10
1.074706727 2.376680750 2.928139705 3.574377652 0.078062625 -0.113745920 0.045641386 -0.004563741
> coef(best_subset_new, which.min(summary_new$cp))
(Intercept)          X1          X2          X3          X5
1.07200775 2.38745596 2.84575641 3.55797426 0.08072292
> coef(best_subset_new, which.min(summary_new$adjr2))
(Intercept)          X3
3.437156 4.828270
> |
```




BIC (lowest) selected 3 variables (X1, X2, X3). This is wrong, because the true model only contains X7. Cp (lowest) selected 5 variables (X1, X2, X3, X5). Again, this is incorrect. Adjusted R² (highest) selected 1 variable (X3) but still wrong.

Hence, best subset selection completely failed to identify X7. It was misled by the strong correlations between X, X2, ..., X10 and the noise.

```

> library(glmnet)
> x.mat.new = model.matrix(Y_new ~ poly(X, 10, raw = TRUE))[, -1]
> y.vec.new = Y_new
> set.seed(1)
> cv.lasso.new = cv.glmnet(x.mat.new, y.vec.new, alpha = 1)
> plot(cv.lasso.new)
> best.lambda.new = cv.lasso.new$lambda.min
> best.lambda.new
[1] 14.13562
> lasso.fit.new = glmnet(x.mat.new, y.vec.new, alpha = 1, lambda = best.lambda.new)
> coef(lasso.fit.new)
11 x 1 sparse Matrix of class "dgCMatrix"
              s0
(Intercept)  1.947394330
poly(X, 10, raw = TRUE)1 .
poly(X, 10, raw = TRUE)2 .
poly(X, 10, raw = TRUE)3 .
poly(X, 10, raw = TRUE)4 .
poly(X, 10, raw = TRUE)5 .
poly(X, 10, raw = TRUE)6 .
poly(X, 10, raw = TRUE)7  7.750481484
poly(X, 10, raw = TRUE)8 .
poly(X, 10, raw = TRUE)9  0.002855533
poly(X, 10, raw = TRUE)10 .

```



Lasso correctly selected X_7 and eliminated all irrelevant predictors. The estimated coefficient of 7.75 is very close to the true value $\beta_7=8$. Hence, Lasso outperformed best subset selection when the true model has only one predictor (X_7).